

MANTRA: A reference dataset for astronomical transient event recognition. Part I: lightcurves

MAURICIO NEIRA,¹ CATALINA GÓMEZ,² DIEGO A. GÓMEZ,¹ JUAN PABLO REYES,¹ MARCELA HERNÁNDEZ HOYOS,¹
PABLO ARBELÁEZ,² AND JAIME E. FORERO-ROMERO³

¹ *Systems and Computing Engineering Department
Universidad de los Andes*

*Cra. 1 No. 18A-10
Bogotá, Colombia*

² *Departamento de Ingeniería Biomédica
Universidad de los Andes*

*Cra. 1 No. 18A-10
Bogotá, Colombia*

³ *Departamento de Física
Universidad de los Andes
Cra. 1 No. 18A-10
Bogotá, Colombia*

Submitted to ApJS

ABSTRACT

We introduce MANTRA, an annotated dataset of 4869 transient and 16940 non-transient object lightcurves built from the Catalina Real Time Transient Survey. We provide public access to this dataset as a plain text file to facilitate standardized quantitative comparison of astronomical transient event recognition algorithms. Some of the classes included in the dataset are: supernovae, cataclismic variables, active galactic nuclei, high proper motion stars, blazars and flares. As a complement to the dataset, we experiment with multiple data pre-processing methods, feature selection techniques and popular machine learning algorithms (Support Vector Machines, Random Forests and Neural Networks). We assess quantitative performance in two classification tasks: binary (transient/non-transient) and eight-class classification. The best performing algorithm in both task is the Random Forest Classifier. It achieves an F1-score of 86.61% in the binary classification and 50.38% in the eight-class classification. For the latter, the class with the highest F1-score are non-transients (87.12%) and the lowest corresponds to flares (11.96%); for supernovae it achieves a value of 50.07%, close to the average across classes. The next release of MANTRA includes images and benchmarks with deep learning models.

Keywords: Astronomy databases (83), Transient detection (1957), Astrostatistics tools (1887)

1. INTRODUCTION

Large scale automatic detection and classification of astronomical transients is happening within surveys such as Pan-STARRS1 (Kaiser 2004), the Palomar Transient Factory (Law et al. 2009), the Catalina Real-Time Transient Survey (Drake et al. 2009), the All-Sky Automated Survey for SuperNovae (Shappee et al. 2014) and the Zwicky Transient Factory (Bellm et al. 2019). Besides the large amount of data, transient classification is hard because the data is usually heterogeneous, unbalanced, sparse, unevenly sampled and with missing information.

These two characteristics (size and heterogeneity) have motivated the application of Machine Learning (ML) algorithms to face this challenge. For instance, Random Forests, MultiLayer Perceptron and K-Nearest Neighbours have been used on lightcurves to classify transients from the Catalina Real Time Transient Survey (D’Isanto et al. 2016); convolutional neural networks have been used as input to automatic vetting algorithms (quick classification of bogus vs. real transients) based on data from the SkyMapper Supernova and Transient Survey and the High cadence Transient Survey (HiTS) (Gieseke et al. 2017; Cabrera-Vives et al. 2017).

The quality of the dataset *is a necessary but insufficient* requirement for the success of these examples and of any other ML implementation. New ML results usually come from groups internal to an observational collaboration because they have the internal know-how (and, sometimes, privileged access) to build training datasets. Consequently, despite having a lot of data available in the public domain, much of it is difficult to access. This difference in data access makes it challenging for the broader astronomical community to rebuild a training dataset, perform comparisons with published results and suggest new algorithms.

Nonetheless, building training datasets has been facilitated by the publication of large databases from different observational projects. Other collaborations have directly published large datasets of *simulated transients* hoping to trigger more involvement from the ML in astronomy community at large to develop new classification algorithms (The PLAsTiCC team et al. 2018). However, at the time of writing, *no dataset for transients has been made accessible to the public in the form of a catalogue based on real data*.

To this end we compile and publish in easy-to-access files a dataset that can be used to train and test different ML algorithms for transient classification. We use public data from the Catalina Real-Time Transient Survey (CRTS) (Drake et al. 2012), an astronomical survey searching transient and highly variable objects as base for the dataset. Effectively, we developed an ETL (extract, transform and load) procedure to extract the data from CRTS, which is designed for a sporadic lookup of a few objects, into a catalogue of thousands of objects that can be used to train ML algorithms and establish benchmarks. Here, in the first paper, we present the lightcurve data. In a second paper we will present a curated imaging dataset from the same survey.

This paper is structured as follows. In Section 2 we present the CRTS and the steps we follow to build the dataset. Then, in Section 3 we describe its main features together with the repository structure gathering the files and Python code to explore it. In Section 4 we show how this dataset can be used to perform tests using ML methods following a similar approach as D’Isanto et al. (2016), and the experiments that we perform. We finalize in Section 5 with a summary of the main features of our dataset and the results of our experiments.

2. THE LIGHTCURVE DATASET

We use public data from the Catalina Real-Time Transient Survey (CRTS) (Drake et al. 2009), an astronomical survey searching transient and highly variable objects. The CRTS covered 33000 squared degrees of

sky and took data since 2007. Three telescopes were used: Mt. Lemmon Survey (MLS), Catalina Sky Survey (CSS), and Siding Spring Survey (SSS). So far, CRTS has discovered more than 15000 transient events. We use data from the CSS telescope, which is an f/1.8 Schmidt telescope located in the Santa Catalina Mountains in Arizona. The telescope is equipped with a 111-megapixel detector, and covered 4000 square degrees per night, with a limiting magnitude of 19.5 in the V band.

Putting together the lightcurves for MANTRA implies cross-matching different files in the legacy CRTS webpage: http://nesssi.cacr.caltech.edu/DataRelease/CRTS-I_transients.html. The photometry is stored in two different kinds of files: `phot` that come from the main photometry database and `orphan` that correspond to transients not associated with the 500 million sources in the main photometry database. There are also `out` files that must be used to link transient IDs to database IDs.

For each one of the 5540 transients reported and classified in the archival webpage <http://nesssi.cacr.caltech.edu/catalina/All.arch.html> we use its transient IDs and its database IDs to look for the lightcurves in the `phot` and `orphan` files. Only 4982 transients can be linked to available data to reconstruct their lightcurves. Furthermore, some of these lightcurves are duplicated, i.e. they had the same number of observations, Modified Julian Date (MJD) and magnitude measurement. We ignore the duplicates to end up with 4869 unique transients with an associated lightcurve. Figure 1 summarizes this process.

The CRTS dataset already provides a classification. The most numerous classes are: supernovae, cataclysmic variable stars, blazars, flares, asteroids, active galactic nuclei, and high-proper-motion stars (HPM). Though most objects in the transient object catalogue belong to a single class, there is some uncertainty in the categorization of some of them. In this case, an interrogation sign is used when a class is not clear e.g. SN? or sometimes multiple possible classes are found for a single event e.g. SN/CV. Table 1 summarizes the number of objects in each class.

To compile the non-transient lightcurves we retrieve sources in the dataset from the CRTS online catalogue by retrieving objects within a 20 arcsecond radius from CRTS detected transients, and removing known transient lightcurves from that set. We end up with 16940 unique non-transient lightcurves. Figure 2 illustrates this process.

Figure 3 shows the number of lightcurves as a function of average magnitude (left panel) and as a function of the number of points in the lightcurve (right panel).

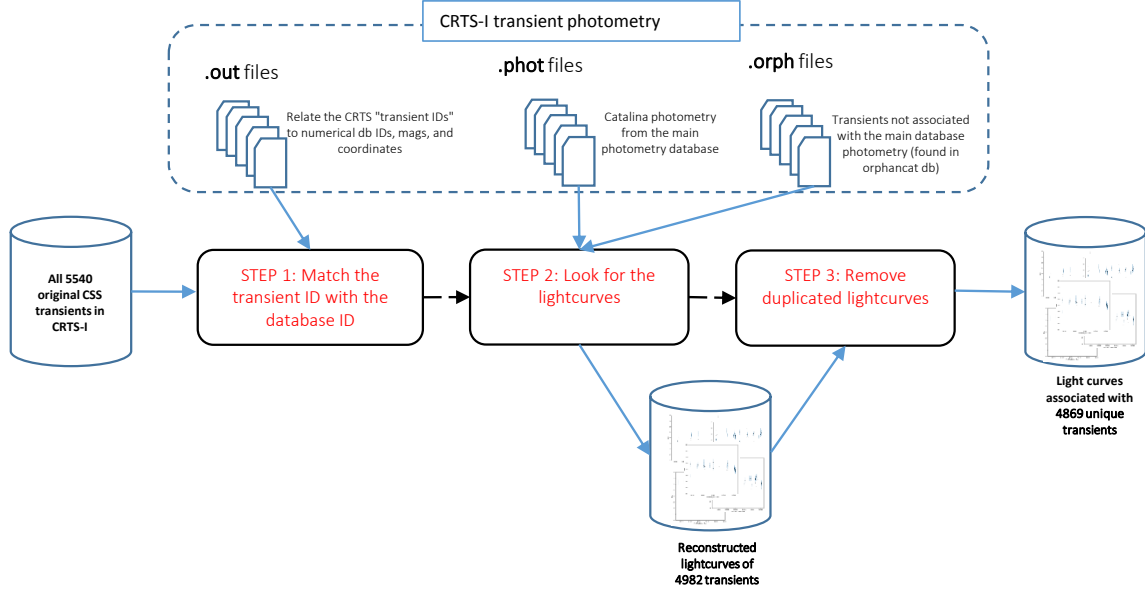


Figure 1. MANTRA Dataset Set Up: Lightcurve compilation for transient classes.

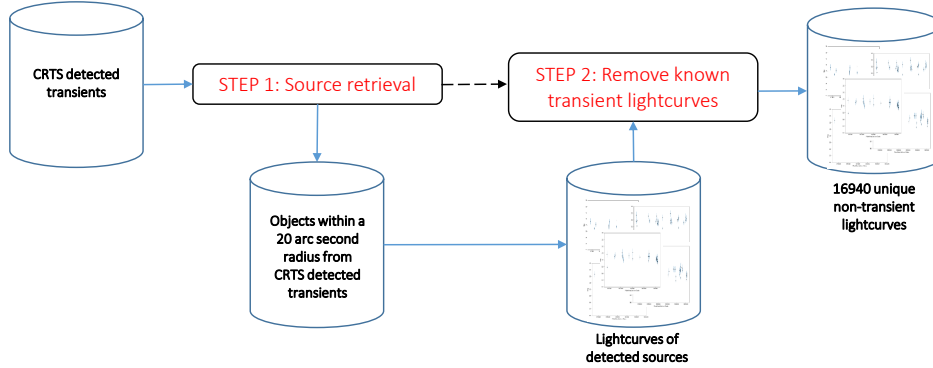


Figure 2. MANTRA Dataset Set Up: Lightcurve compilation for non-transients.

We show separately the whole data set and three most representative classes: supernova, cataclysmic variables and active galactic nuclei. For these four sets, the median magnitude is in the range 18 – 20. The number of points in the lightcurve has a larger variability. The median for all the curves is close to 30, while for SN, CV and AGN it is close to 15, 50 and 180, respectively. We provide sample lightcurves of the most represented transient classes and non-transient sources in Figure 4 and Figure 5, respectively. The brightness evolution of non-transient sources is more stable over time, while transient objects present non-periodical changes at different time scales.

The challenges in the classification of lightcurves include the inherent nature of transient events, which is reflected in different brightness behaviors, their evolution over time, and the nonuniform sampling of observations at sequential dates. Besides, there is a large class imbalance to localize transient events, and perform their subsequent classification.

2.1. Classification Tasks

We study two classification tasks on the MANTRA dataset:

- **Binary Classification.** Using a balanced number of events from both classes in order to investigate




Figure 3. Cumulative number of lightcurves (expressed as a fraction) as a function of average magnitude (left) and number of data points in the lightcurve (right). This includes information for the three most representative classes (SN, CV, AGN) and the whole database (ALL).

Class	Object Count
SN	1723
CV	988
HPM	640
AGN	446
SN?	319
Blazar	243
Unknown	228
Flare	219
AGN?	138
CV?	77

Table 1. Top 10 transient classes in the CRTS with their respective number of lightcurves.

the capability of distinguishing between transient and non-transient sources.

- **8-Class Classification.** Using the unbalanced number of objects across classes to perform a classification into the following categories: AGN, Blazar, CV, Flare, HPM, Other, SN and Non-Transient.

We evaluate both tasks using the metrics of a detection problem. For each class in the testing set, we report the maximum F1-Score that is defined as the harmonic mean of precision and recall. We construct Precision-Recall (PR) curves by setting different thresholds on the output probabilities of belonging to each class.

3. REPOSITORY DESCRIPTION

The repository contains the lightcurves and a Jupyter notebook to reproduce some of the Figures and Tables in this paper. The repository can be found in <https://github.com/MachineLearningUniandes/MANTRA>. To date the repository has two main folders:

- **data/lightcurves:** contains three text files in CSV format the transient lightcurves (`transient_lightcurves.csv`), the labels for the transients (`transient_labels.csv`) and the lightcurves for non-transient objects (`nontransient_lightcurves.csv`). The first two files can be linked by unique transient IDs and provided in the CRTS database.
- **nb-explore:** includes a jupyter notebook (`explore_light_curves.ipynb`) with examples on how to read and plot transient and non-transient lightcurves, extract the statistics in Table 1 and prepare the summary statistics in Figure 3. Additional python files (`features.py`, `helpers.py` and `inputs.py`) allow to read and perform simple operations on the CSV data files.

4. MACHINE LEARNING METHODS

Here we test baseline algorithms on the MANTRA dataset that can be used as a reference for future work. Figure 6 shows an overview of our transient classification framework. The main steps include data processing, feature extraction and classification.

4.1. Preprocessing and Feature Extraction

We do not input directly the annotated lightcurves to the ML algorithms. We perform a preprocessing stage as follows. First, we discard lightcurves with less than 10 data points observations as they may not contain enough information to be classified correctly.

Given that the number of lightcurves per class is imbalanced, in order to have the same number of instances for each class, we implement an oversampling step by artificially generating multiple mock lightcurves from an observed one. We generate a slightly different lightcurve

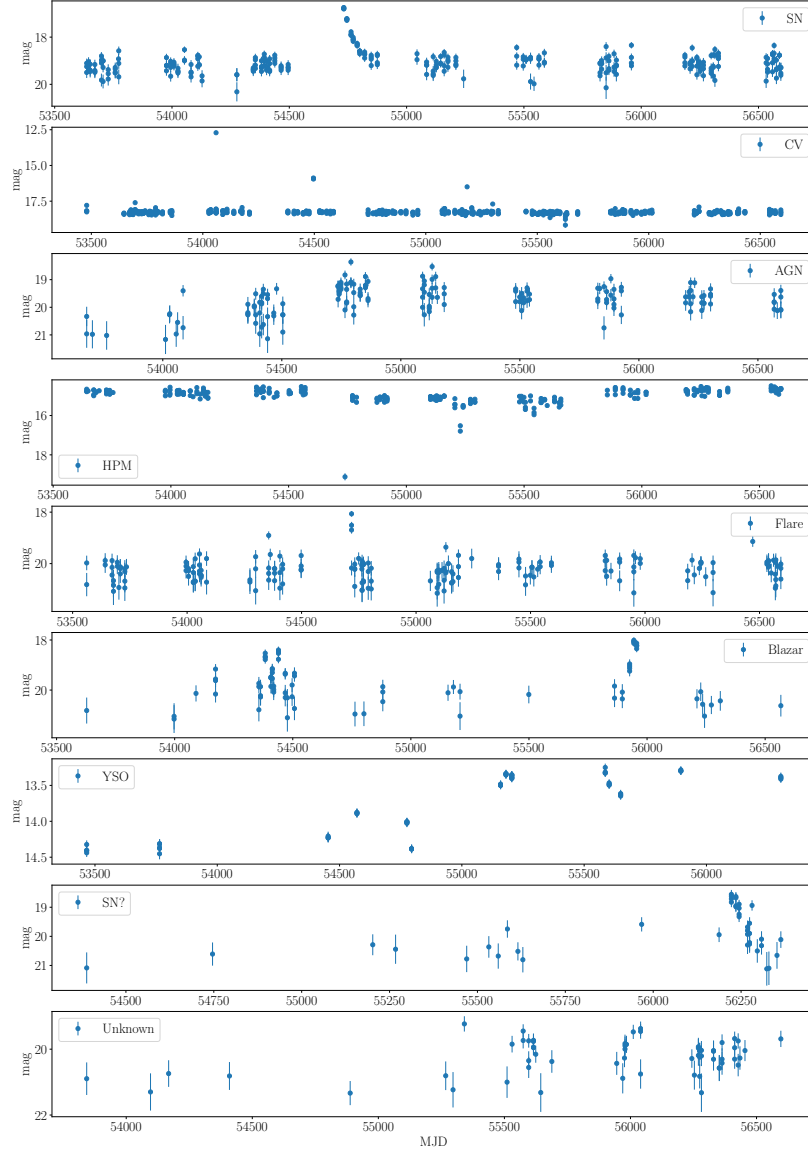


Figure 4. Randomly selected lightcurves for the most represented transient classes as compiled in MANTRA. The class of each sample is within the legend box.

from the observed lightcurve and then sample the observed magnitude from a Gaussian distribution centered on the observational apparent magnitude with the magnitude’s error as the standard deviation. It is important to note that *the oversampling was only done on the train set. The test set was left untouched.*

Finally, we compute a standard set of features for each lightcurve. These features are scalars derived from statistical and model-specific fitting techniques. The first features (moment-based, magnitude-based and percentile-based) were formally introduced in Richards et al. (2011), and have been used in other studies (Lochner et al. 2016; D’Isanto et al. 2016). We extend that list to include another set (polynomial fitting-based

features). At the end of this process we normalize the features to have zero mean and unit variance.

These groups of features are:

1. Moment-based features, which use the magnitude for each lightcurve.
 - **beyond1std:** Percentage of observations which are over or under one standard deviation from the weighted average. Each weight is calculated as the inverse of the corresponding observation’s photometric error.
 - **kurtosis:** The fourth moment of the data distribution.

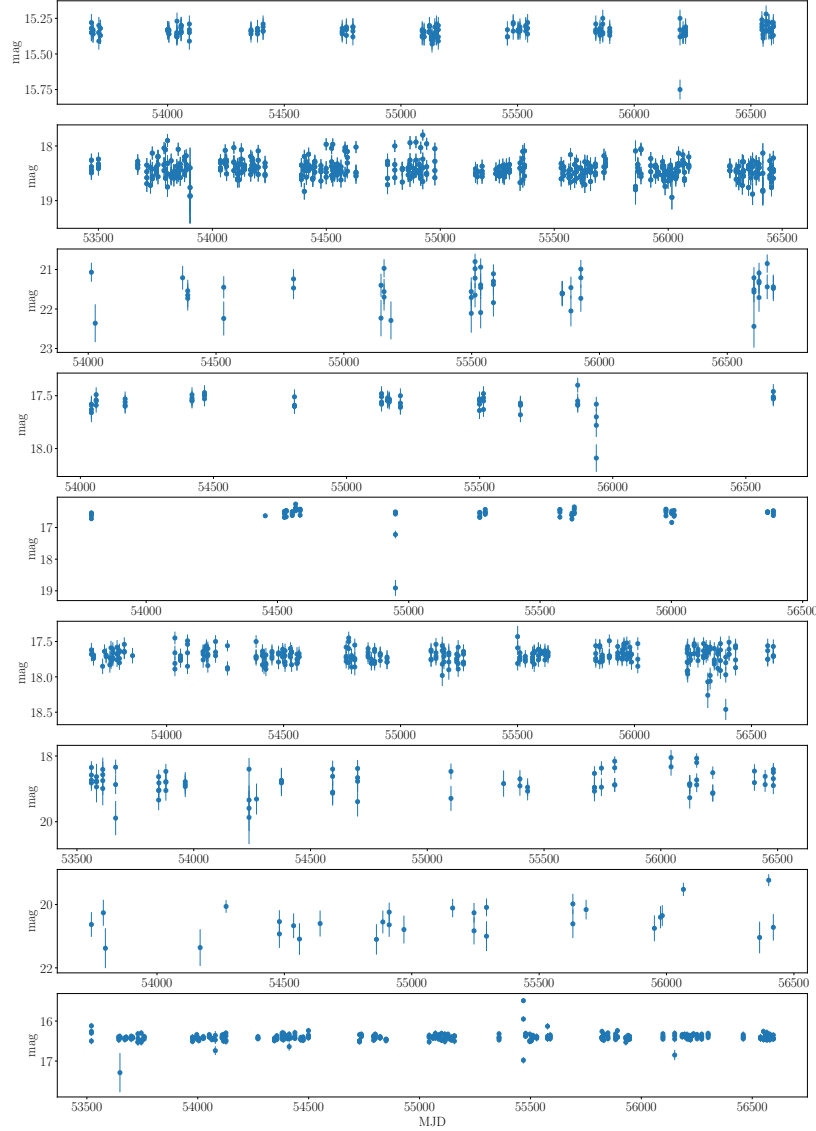


Figure 5. Randomly selected lightcurves for non-transient sources retrieved for MANTRA.

- **skew:** Skewness. Third moment of the data distribution.
- **sk:** Small sample kurtosis.
- **std::** The standard deviation.
- **stetson_j:** The Welch-Stetson J variability index (Stetson 1996). A robust standard deviation.
- **stetson_k:** The Welch-Stetson K variability index (Stetson 1996). A robust kurtosis measure.
- **max_slope:** Maximum absolute slope between two consecutive observations.
- **mad:** The median of the difference between magnitudes and the median magnitude.
- **mbrp:** The percentage of points within 10% of the median magnitude.
- **pst:** Percentage of all pairs of consecutive magnitude measurements that have positive slope.
- **pst_last30:** Percentage of the last 30 pairs of consecutive magnitudes that have a positive slope, minus percentage of the last 30 pairs of consecutive magnitudes with a negative slope.

2. Features based on the magnitudes.

- **amp:** The difference between the maximum and minimum magnitudes.

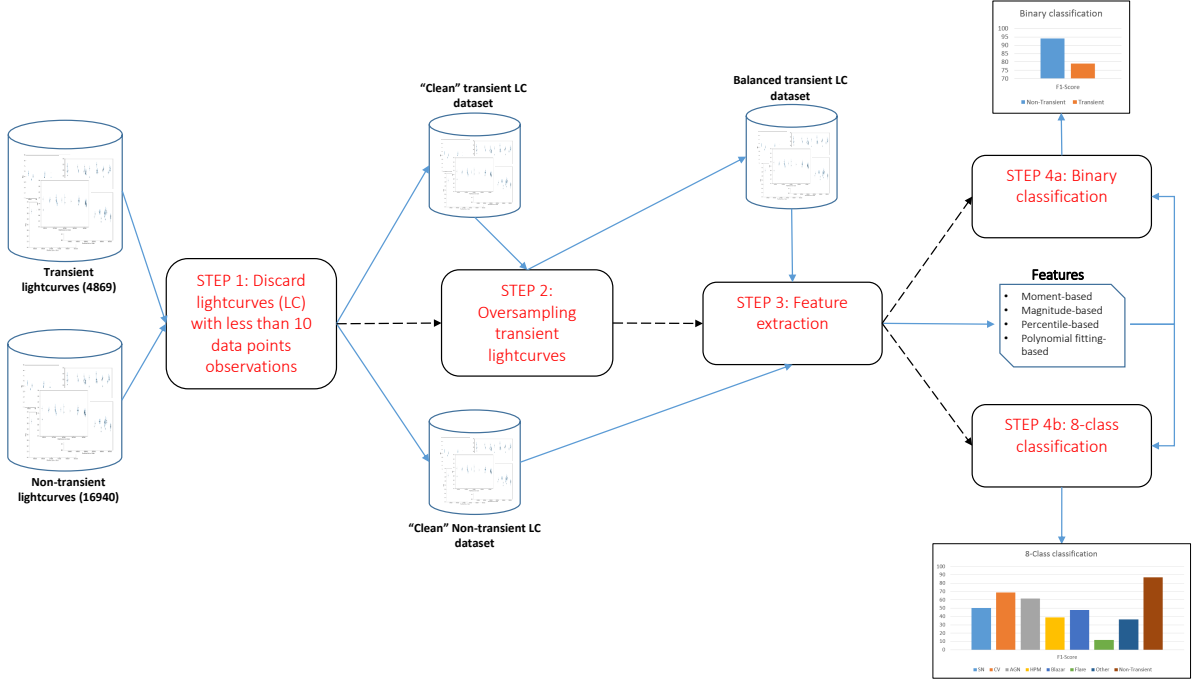


Figure 6. Overview of the Machine Learning process on the MANTRA dataset for the binary and 8-class classification tasks. We take the raw lightcurves as input, preprocess the data (step 1) and balance the classes for the training phase (step 2). We extract the features (step 3) that are feed into the ML algorithms that perform the classification task (step 4).

3. Percentile-based features, which use the sorted flux distribution for each source. The flux is computed as $F = 10^{0.4\text{mag}}$. We define $F_{n,m}$ as the difference between the m -th and n -th flux percentiles.

- **p_amp**: Largest percentage difference between the absolute maximum magnitude and the median.
- **pdfp**: Ratio between $F_{5,95}$ and the median flux.
- **fpr20**: Ratio $F_{40,60}/F_{5,95}$
- **fpr35**: Ratio $F_{32.5,67.5}/F_{5,95}$
- **fpr50**: Ratio $F_{25,75}/F_{5,95}$
- **fpr65**: Ratio $F_{17.5,82.5}/F_{5,95}$
- **fpr80**: Ratio $F_{10,90}/F_{5,95}$

4. Polynomial Fitting-based features, which are the coefficients of multi-level terms in a polynomial curve fitting. This is a new set of features proposed in this paper. **Polyn_Tm** indicates the coefficient of the term of order m in a fit to a polynomial of order n .

- **Poly1_T1**.
- **Poly2_T1**.

- **Poly2_T2**.
- **Poly3_T1**.
- **Poly3_T2**.
- **Poly3_T3**.
- **Poly4_T1**.
- **Poly4_T2**.
- **Poly4_T3**.
- **Poly4_T4**.

4.2. ML algorithms

We conduct experiments with three widely used families of supervised classification algorithms (Bloom et al. 2012; D’Isanto et al. 2016): Neural Networks (NNs), Random Forests (RFs) and Support Vector Machines (SVMs).

These algorithms are popular in published studies and are efficient for low dimensional feature datasets as is our case. We use sklearn (Pedregosa et al. 2011) Python’s implementation of random forests and support vector machines. Details on the inner workings of these machine learning models can be found in Hastie et al. (2016).

We use the pytorch library for python for the development of the linear Neural Networks. It consists of a series of fully connected layers that map the features

Case	Classifier	Precision	Recall	F1-score
Binary	RF	86.61	86.61	86.61
	SVM	82.61	82.525	82.57
	NN	71.84	73.19	72.51
8 Class	RF	46.25	63.59	50.38
	SVM	32.94	55.22	36.60
	NN	36.0	55.19	38.59

Table 2. Average precision, recall and F1-score across all classes for each algorithm and classification task. Best results per metric per classification task are in bold.

to the corresponding number of classes. At each layer, a 1d batch normalization is implemented followed by a relu activation function. The final layer calls a softmax activation function to transform the numerical values to class probabilities.

The hyperparameters explored for each algorithm are the following.

- Neural Networks:
 - Learning Rate: $\{0.1, 0.01, 0.001, 0.0001\}$
 - Hidden Layer Sizes: Single, double or triple layers with 500 nodes each.
- Random Forest:
 - Number of Estimators: 200 or 700.
 - Number of features considered: Square root or \log_2 of the total number of features.
- Support Vector Machines:
 - Kernels: Radial Basis Function (RBF), linear or sigmoid.
 - Kernel Coefficient (γ): $\{0.125, 2, 32\}$
 - Error Penalty (C): $\{0.125, 2, 32\}$

4.3. Validation

We split the input lightcurves into training and testing in a 75 : 25 ratio respectively, class by class. For the random forests and the SVM, we use a grid search over the hyperparameter combinations with a 2-fold cross-validation over the training set to determine the best hyperparameters. For the neural networks, at each epoch, the network is evaluated on the test data.

4.4. Results

Table 2 shows the average class precision, recall and F1-measure for each of the classification tasks and algorithms listed above.

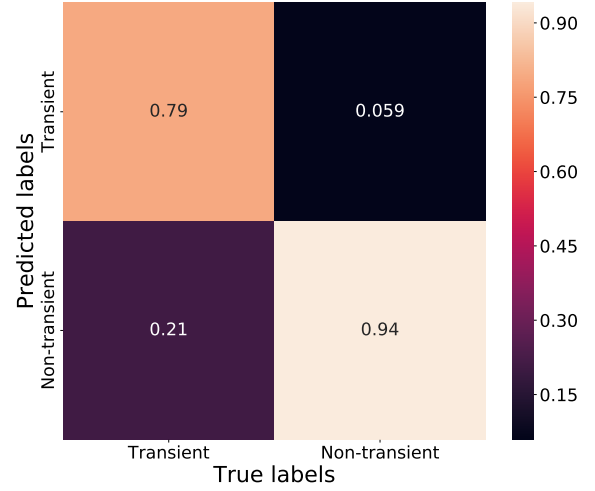


Figure 7. Confusion Matrix for the best performing model in the Binary task. Rows represent prediction and columns the ground truth.

4.4.1. Binary Classification

The best algorithm in this task is RFs with an average F1-Score of 86.61%. SVMs are the second best-performing model with a F1-Score of 82.57%. Changing the number of features does not significantly affect the score. NNs are ranked third with an F1-Score of 72.51%.

Figure 7 shows the confusion matrix of the best performing algorithm. These results suggest that in an imbalanced set up, non-transient sources are better classified while transients are more difficult, showing a difference of about 15 points in the percentage of correct classifications. This difference in performance could be attributed to the intra-class variation within the transients.

4.4.2. Eight-Class Classification

For this task, RF is again the best classifier. The best F1-Score is 50.38%. NNs are the second best model with an F1-Score of 36.60%, while SVMs are the worst-performing model only achieving an average F1-Score of 38.59%. Table 3 summarizes the results for individual classes and Figure 8 presents the confusion matrix for the RF.

The two classes with highest F1-Score are non-transient (87.12%) and CV (68.77%). The recall decreases for the non-transient class in comparison to the binary experiment, meaning that the algorithm misclassified some instances that belong to non-transient class among transient classes. However, transient sources are not commonly confused with non-transient ones. The




Figure 8. Confusion Matrix for the best performing model in the 8-class task. The classes follow the abbreviations in Table 3. Rows represent the predictions, and columns the ground truth.

worst performing classes are Flare, Other and HPM, with F1-Scores in the range 11% - 40%. It is worth noting that the less frequent classes present a lower performance, such as Flare and HPM. Even though the most frequent classes are more easily identified, the "other" type class has a low F1-score due to the diverse nature of sources assigned to this category.

5. CONCLUSIONS

The scope of forthcoming large astronomical synoptic surveys motivates the development and exploration of automatized ways to detect transient sources. Such developments require observational datasets to train and test new algorithms. Making these datasets public and easy to access has the potential to open the field to a larger number of contributors. With these goals in mind, in this paper we presented a compilation based on data from the Catalina Real-Time Transient Survey. The dataset has 4869 transient and 16940 non-transient

Class	Precision	Recall	F1-score	Cover
SN	48.82	51.39	50.07	323
CV	66.96	70.69	68.77	215
AGN	48.14	85.84	61.69	106
HPM	25.19	86.84	39.05	76
Blazar	46.77	49.15	47.93	59
Flare	7.00	41.17	11.96	51
Other	31.11	44.01	36.46	234
Non-Tr.	96.06	79.69	87.12	3798
avg/total	46.25	63.60	50.38	4862

Table 3. Precision, Recall and F1-Score for the 8-Class Classification Task.

lightcurves. The dataset is publicly available at <https://github.com/MachineLearningUniandes/MANTRA>.

We illustrated how to use this database by extracting characteristic features to use them as input to train three different machine learning algorithms (Random Forests, Neural Networks and Support Vector Machines) for clas-

sification tasks. The features extracted from lightcurves were either statistical descriptors of the observations, or polynomial curve fitting coefficients applied to the lightcurves. Overall, the best classifier for all tasks was the Random Forest. Neural Networks showed the worst performance in the binary classification task and the second best in the 8-class classification task.

Certainly other classification algorithms could be used on the MANTRA lightcurves. Our purpose here is not a thorough analysis of machine learning algorithms. Our focus is in the compilation of the data into a format that could easily be used for different projects that could be addressed with MANTRA. For instance the supernovae detection problem could be studied as a classification problem between SN and the other classes, including non-transients. Considering extremely unbalanced classes could also be addressed with the dataset we are presenting by simply using less samples for the SN class and keeping a large numbers for the non-transient class.

In a second paper we will present the second part of MANTRA. It corresponds to almost one million images

from the CRTS. These will be tested using state-of-the-art deep learning techniques for transient classification.

ACKNOWLEDGEMENTS

We thank Andrew Drake for sharing with us the CRTS Transient dataset used in this project. We acknowledge funding from Universidad de los Andes in the call for project finalization. We also thank contributors and collaborators of the SciKit-Learn, Jupiter Notebooks and Pandas' Python libraries.

CRTS and CSDR2 are supported by the U.S. National Science Foundation under grant NSF grants AST-1313422, AST-1413600, and AST-1518308. The CSS survey is funded by the National Aeronautics and Space Administration under Grant No. NNG05GF22G issued through the Science Mission Directorate Near-Earth Objects Observations Program.

REFERENCES

- Bellm, E. C., Kulkarni, S. R., Graham, M. J., et al. 2019, Publications of the Astronomical Society of the Pacific, 131, 018002
- Bloom, J. S., Richards, J. W., Nugent, P. E., et al. 2012, Publications of the Astronomical Society of the Pacific, 124, 1175
- Cabrera-Vives, G., Reyes, I., Förster, F., Estévez, P. A., & Maureira, J.-C. 2017, ApJ, 836, 97
- D'Isanto, A., Caviuoti, S., Brescia, M., et al. 2016, MNRAS, 457, 3119
- Drake, A. J., Djorgovski, S. G., Mahabal, A., et al. 2009, ApJ, 696, 870
- Drake, A. J., Djorgovski, S. G., Mahabal, A., et al. 2012, in IAU Symposium, Vol. 285, New Horizons in Time Domain Astronomy, ed. E. Griffin, R. Hanisch, & R. Seaman, 306–308
- Gieseke, F., Bloemen, S., van den Bogaard, C., et al. 2017, MNRAS, 472, 3101
- Hastie, T., Tibshirani, R., & Friedman, J. 2016, The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Second Edition (Springer Series in Statistics) (Springer)
- Kaiser, N. 2004, in Society of Photo-Optical Instrumentation Engineers (SPIE) Conference Series, Vol. 5489, Ground-based Telescopes, ed. J. Oschmann, Jacobus M., 11–22
- Law, N. M., Kulkarni, S. R., Dekany, R. G., et al. 2009, PASP, 121, 1395
- Lochner, M., McEwen, J. D., Peiris, H. V., Lahav, O., & Winter, M. K. 2016, ApJS, 225, 31
- Pedregosa, F., Varoquaux, G., Gramfort, A., et al. 2011, Journal of machine learning research, 12, 2825
- Richards, J. W., Starr, D. L., Butler, N. R., et al. 2011, ApJ, 733, 10
- Shappee, B. J., Prieto, J. L., Grupe, D., et al. 2014, ApJ, 788, 48
- Stetson, P. B. 1996, pasp, 108, 851
- The PLAsTiCC team, Allam, Tarek, J., Bahmanyar, A., et al. 2018, arXiv e-prints, arXiv:1810.00001